


Ruchita Paithankar
CS583
Homework 1

- **split(img)**
 1. figuring this out was the most straightforward thing after the merge(r,g,b) function.
 2. It is basically splitting the array depth-wise into 3 channels – R, B, G and squeezing them into an 1D array.
- **create_gaussian_kernel(size,sigma=1.0)**
 1. To compute a 2D gaussian kernel I iterated the loops from -size/2 to size/2 so that the center x,y is 0,0.
 2. Then I used the given formula for calculating Gaussian kernel.
 3. Tried normalization by summing the values of 'rv'.
 4. I got a little stuck here as I was not clear on the concepts at first. I referred to a few concepts from the week 2 slides.


Gaussian filters



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- Like in this image the professor talked about the value of sigma adding up to 1.
5. Initially I forgot to divide the kernel by the sum of its values.

Image filtering



$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$

Credit: S. SeI

Again, I had to brush up some things conceptually. This box filter slide from week 2 is what I referred to.

- **convolve_pixel(img, kernel, i, j)**

1. this was where I got stuck the most. Basically, figuring out the condition for when the kernel would land outside the image box, and then the answer I was initially 18.6 instead of 18.4

```

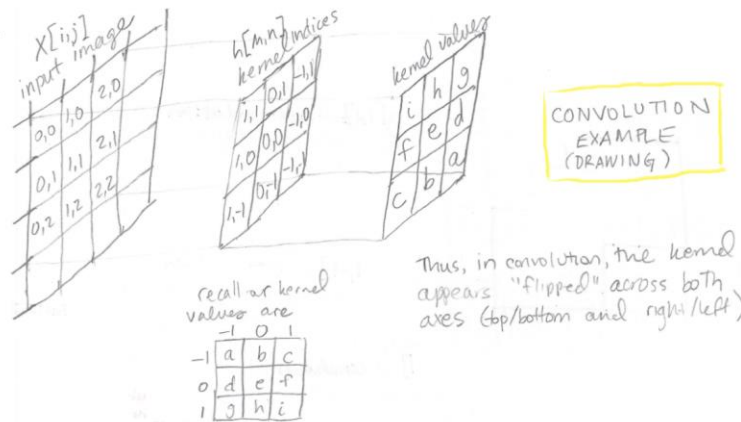
FAILED (errors=1)
(base) ruchitapaithankar@Ruchitas-MacBook-Air hw1_kit(1) % python3 hw1_test.py Homework1Test.test_convolve_pixel -v
test_convolve_pixel (__main__.Homework1Test) ... FAIL

=====
FAIL: test_convolve_pixel (__main__.Homework1Test)
=====
Traceback (most recent call last):
  File "hw1_test.py", line 48, in test_convolve_pixel
    self.assertEqual(result, 18.4, places=4,
AssertionError: 18.6 != 18.4 within 4 places (0.20000038146972798 difference) : The correct result is computed.

=====
Ran 1 test in 0.001s

```

2. I went through the slack chat and someone had put in a comment like, it correlates to 18.6 but it convolutes to 18.4. That is when I realized where I went wrong.



This is a picture from a website I landed on when I looked for convolution vs correlation. I tried transposing first on paper, but that won't give me the desired matrix. I tried transpose + mirror image, that didn't work too.

Convolution vs. image filtering

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

Filtering (also known as cross correlation)

$$h[m,n] = \sum_{k,l} f[k,l] I[m-k,n-l]$$

Convolution

Remembered the professor mentioning that x changes to y and y to x. So, I thought of rotating the matrix 90 degree and then rotating it once more to complete a 180 degree rotation, gave me the kernel I needed.

3. Taking this reverse kernel row by row and column by column on the image, multiplying elements by each other and then summing up the result.
4. Again, I used $-size/2$ to $size/2$ of the kernel to iterate through the pixels.


- **convolve(img, kernel):**

1. convoluting the image pixel by pixel (x,y) with the kernel from the previous function.

```
(base) ruchitapaithankar@Ruchitas-MacBook-Air hw1_kit(1) % python3 hw1.py --k 3 --sigma 1 example-input1.jpg example-output1.jpg
INFO: Loading input image example-input1.jpg
INFO: Splitting it into 3 channels
INFO: Computing a gaussian kernel with size 3 and sigma 1.000000
INFO: Convolving the first channel
INFO: Convolving the second channel
INFO: Convolving the third channel
INFO: Merging results
INFO: Saving result to example-output1.jpg
(base) ruchitapaithankar@Ruchitas-MacBook-Air hw1_kit(1) % python3 hw1.py --k 7 --sigma 2 example-input1.jpg example-output1.jpg
INFO: Loading input image example-input1.jpg
INFO: Splitting it into 3 channels
INFO: Computing a gaussian kernel with size 7 and sigma 2.000000
INFO: Convolving the first channel
INFO: Convolving the second channel
INFO: Convolving the third channel
INFO: Merging results
INFO: Saving result to example-output1.jpg
```

Tried running the file for the example-input1 which $k=3$ and $\sigma=1$ and didn't see much change. So tried, $k=7$ and $\sigma=2$ and noticed slight blur on the edges of some sharp lines.

Below is the side by side comparison of the two example images I tested on. The first example doesn't show much difference. The second example show a little more noticeable difference.



```
INFO: Saving result to example-output.jpg
(base) ruchitapaithankar@Ruchitas-MacBook-Air hw1_kit(1) % python3 hw1.py --k 7 --sigma 2 example-input2.jpg example-output2.jpg
INFO: Loading input image example-input2.jpg
INFO: Splitting it into 3 channels
INFO: Computing a gaussian kernel with size 7 and sigma 2.000000
INFO: Convolving the first channel
INFO: Convolving the second channel
INFO: Convolving the third channel
INFO: Merging results
INFO: Saving result to example-output2.jpg
(base) ruchitapaithankar@Ruchitas-MacBook-Air hw1_kit(1) % python3 hw1_test.py -v
test_convolve (__main__.Homework1Test) ... ok
test_convolve_pixel (__main__.Homework1Test) ... ok
test_create_gaussian_kernel (__main__.Homework1Test) ... ok
test_merge_image (__main__.Homework1Test) ... ok
test_split_image (__main__.Homework1Test) ... ok
```

